

# TSA Final Project

[https://github.com/vivianzzzzz/ZhangXiaGupta\\_ENV797\\_TSA\\_FinalProject](https://github.com/vivianzzzzz/ZhangXiaGupta_ENV797_TSA_FinalProject)

Xiyue Zhang, Haochong Xia, Ayush Gupta

2024-04-26

```
library(tidyverse)
library(openintro)
library(readxl)
library(lubridate)
library(dplyr)
library(forecast)
library(smooth)
library(zoo)
```

## 1.Data preprocessing

```
bitcoin_raw <- read.csv("data/bitcoin.csv",header=TRUE)

bitcoin <- bitcoin_raw[nrow(bitcoin_raw):1, ] %>%
  mutate(date = as.Date(Date, format="%m/%d/%Y")) %>%
  filter(!is.na('Change')) %>%
  select(date, Price, Change)

bitcoin$Change <- na.locf(bitcoin$Change, na.rm = FALSE)
bitcoin$Change <- gsub("%", "", bitcoin$Change)
bitcoin$Change = as.numeric(as.character(bitcoin$Change))

# Using median to avoid the influence of other potential outliers
median_value <- median(bitcoin$Change[bitcoin$Change != 0], na.rm = TRUE)
bitcoin$Change[bitcoin$Change == 0] <- median_value

# Read the dataset
data <- read.csv("data/DCOILBRENTIU.csv")

# Convert DATE column to Date format
data$DATE <- as.Date(data$DATE)

# Create a sequence of dates from the start date to the end date
start_date <- min(data$DATE)
end_date <- max(data$DATE)
dates <- seq(start_date, end_date, by = "day")

# Filter out Fridays
fridays <- filter(data, weekdays(DATE) == "Friday")
```

```

# Create a dataframe for Saturdays and Sundays
weekend_data <- data.frame(
  DATE = c(fridays$DATE + 1, fridays$DATE + 2),
  DCOILBRETEU = rep(fridays$DCOILBRETEU, each = 1)
)

# Combine the original data with the new weekend data
final_data <- bind_rows(data, weekend_data)

# Sort the final dataset by DATE
final_data <- final_data[order(final_data$DATE), ]

# Reset index
final_data <- final_data %>%
  arrange(DATE) %>%
  mutate(index = row_number()) %>%
  select(index, everything())

# View the modified dataset
#final_data

# Fill missing values with zeros
# final_data_filled <- final_data %>%
#   mutate(DCOILBRETEU = ifelse(is.na(DCOILBRETEU), 0, DCOILBRETEU))

# Convert "." to NA in the DCOILBRETEU column
final_data_filled <- final_data %>%
  mutate(DCOILBRETEU = ifelse(DCOILBRETEU == ".", NA, DCOILBRETEU))

# Replace NA values with the value right before it
final_data_filled <- final_data_filled %>%
  fill(DCOILBRETEU)

# Convert string numbers to numerical values
final_data_filled$DCOILBRETEU <- as.numeric(final_data_filled$DCOILBRETEU)

# View the modified dataframe
#final_data_filled

```

## 1.1 Train test splits & Creating ts and msts

```

# Only use the data after 2017-01-01
bitcoin <- bitcoin %>% filter(date >= ymd("2017-01-01"))

# Create a dummy variable to account for the covid effect. dummy = 1 from march 15 2020 to may 1st 2021
bitcoin <- bitcoin %>% mutate(covid = ifelse(date >= ymd("2020-03-15") & date <= ymd("2021-05-01"), 1, 0))

# Set the proportion for train(90%), test datasets(10%)
train_prop <- 0.9

# Set the number of observations for train, test datasets

```

```

n_train <- floor(nrow(bitcoin) * train_prop)
n_val <- nrow(bitcoin) - n_train

# Split the data into train and test datasets
train_bitcoin <- bitcoin %>% filter(date >= ymd("2017-01-01") & date <= ymd("2023-01-01"))
test_bitcoin <- bitcoin %>% filter(date >= ymd("2023-07-18") )

head(train_bitcoin)

##           date      Price Change covid
## 1 2017-01-01      995.4    3.33      0
## 2 2017-01-02 1,017.00    2.17      0
## 3 2017-01-03 1,033.30    1.60      0
## 4 2017-01-04 1,135.40    9.88      0
## 5 2017-01-05    989.3  -12.86      0
## 6 2017-01-06    886.2  -10.43      0

head(test_bitcoin)

##           date      Price Change covid
## 1 2023-07-18 29,866.80   -0.91      0
## 2 2023-07-19 29,909.70    0.14      0
## 3 2023-07-20 29,801.00   -0.36      0
## 4 2023-07-21 29,903.10    0.34      0
## 5 2023-07-22 29,788.90   -0.38      0
## 6 2023-07-23 30,085.90    1.00      0

ts_bitcoin <- ts(train_bitcoin$Change, frequency = 365.25, start = c(2017,01,01))
ts_bitcoin_test <- ts(test_bitcoin$Change, frequency = 365.25, start = c(2023,07,18))
msts_bitcoin <- msts(train_bitcoin$Change,
                     seasonal.periods = c( 91.25,365.25),
                     start=c(2017,01,01))
#Creating time series with seasonal pattern (quarterly, daily)
msts_bitcoin <- msts(train_bitcoin$Change,
                     seasonal.periods = c( 91.25,365.25),
                     start=c(2017,01,01))
msts_bitcoin_test <- msts(test_bitcoin$Change,
                          seasonal.periods = c( 91.25,365.25),
                          start=c(2023,07,18))

# Set the proportion for train, test datasets
train_prop <- 0.9

# Set the number of observations for train, test datasets
n_train <- floor(nrow(final_data_filled) * train_prop)
n_val <- nrow(final_data_filled) - n_train

# Split the data into train and test datasets
train_oil <- final_data_filled %>% filter(DATE >= ymd("2017-01-01") & DATE <= ymd("2023-01-01"))
test_oil <- final_data_filled %>% filter(DATE >= ymd("2023-07-18") )
#Creating time series
ts_oil <- ts(train_oil$DCOILBRENTU, frequency = 365.25, start = c(2017,01,01))
ts_oil_test <- ts(test_oil$DCOILBRENTU, frequency = 365.25, start = c(2023,07,18))
msts_oil <- msts(train_oil$DCOILBRENTU,
                 seasonal.periods = c( 91.25,365.25),

```

```

start=c(2017,01,01))
msts_oil_test <- msts(test_oil$DCOILBRENTU,
                      seasonal.periods =c( 91.25,365.25),
                      start=c(2023,07,18))


```

## 2. Plot time series and ACF and PACF

```

#plot the time series
ts_bitcoin %>% mstl() %>%
  autoplot()

```




```


# Generate ACF and PACF plots
par(mfrow=c(1,2))
acf(ts_bitcoin, lag.max = 100, main= "ACF plot of bitcoin")
pacf(ts_bitcoin, lag.max = 100, main="PACF plot of bitcoin")

```


ACF plot of bitcoin



PACF plot of bitcoin




```
#plot the time series
ts_oil %>% mstl() %>%
  autoplot()
```




```
# Generate ACF and PACF plots
```

```
par(mfrow=c(1,2))  
acf(ts_oil, lag.max = 100, main= "ACF plot of Crude Oil")  
pacf(ts_oil, lag.max = 100, main="PACF plot of Crude Oil")
```

**ACF plot of Crude Oil**



**PACF plot of Crude Oil**




### 3. Models fit


#### 3.1 STL + ETS model

```
# STL + ETS model  
ETS_fit <- stlf(msts_bitcoin, h=266)  
  
plot(ETS_fit)
```

## Forecasts from STL + ETS(A,N,N)



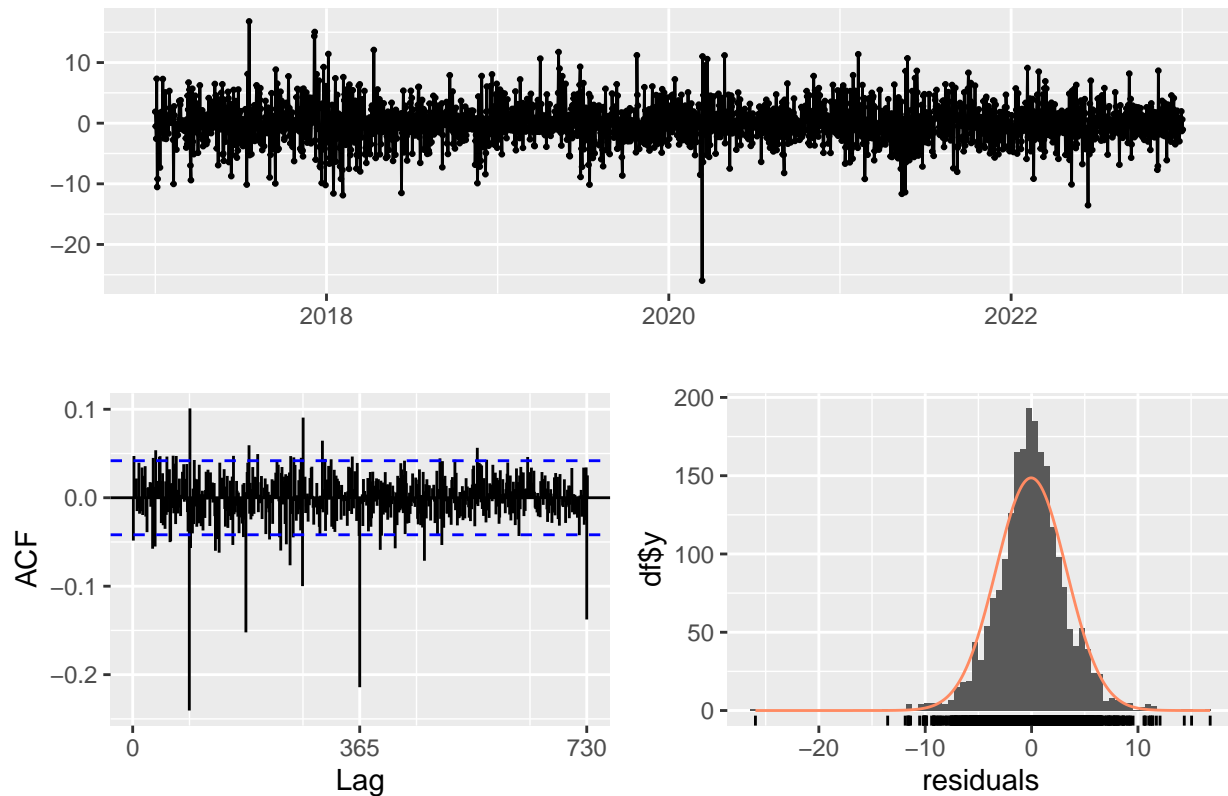
```
autoplot(msts_bitcoin_test) + autolayer(ETS_fit, series="ETS",PI=FALSE)
```



```
ETS_scores <- accuracy(ETS_fit$mean,msts_bitcoin_test)
#print(ETS_scores)

checkresiduals(ETS_fit)
```

### Residuals from STL + ETS(A,N,N)




```
##
##  Ljung-Box test
##
## data:  Residuals from STL +  ETS(A,N,N)
## Q* = 1008.2, df = 438, p-value < 2.2e-16
##
## Model df: 0.   Total lags used: 438
```

### 3.2 TBATS model


```
# TBATS model
TBATS_fit <- tbats(msts_bitcoin)
TBATS_forecast <- forecast(TBATS_fit, h=266)
plot(TBATS_forecast)
```



## Forecasts from TBATS(1, {0,0}, -, {<91.25,6>, <365.25,7>})

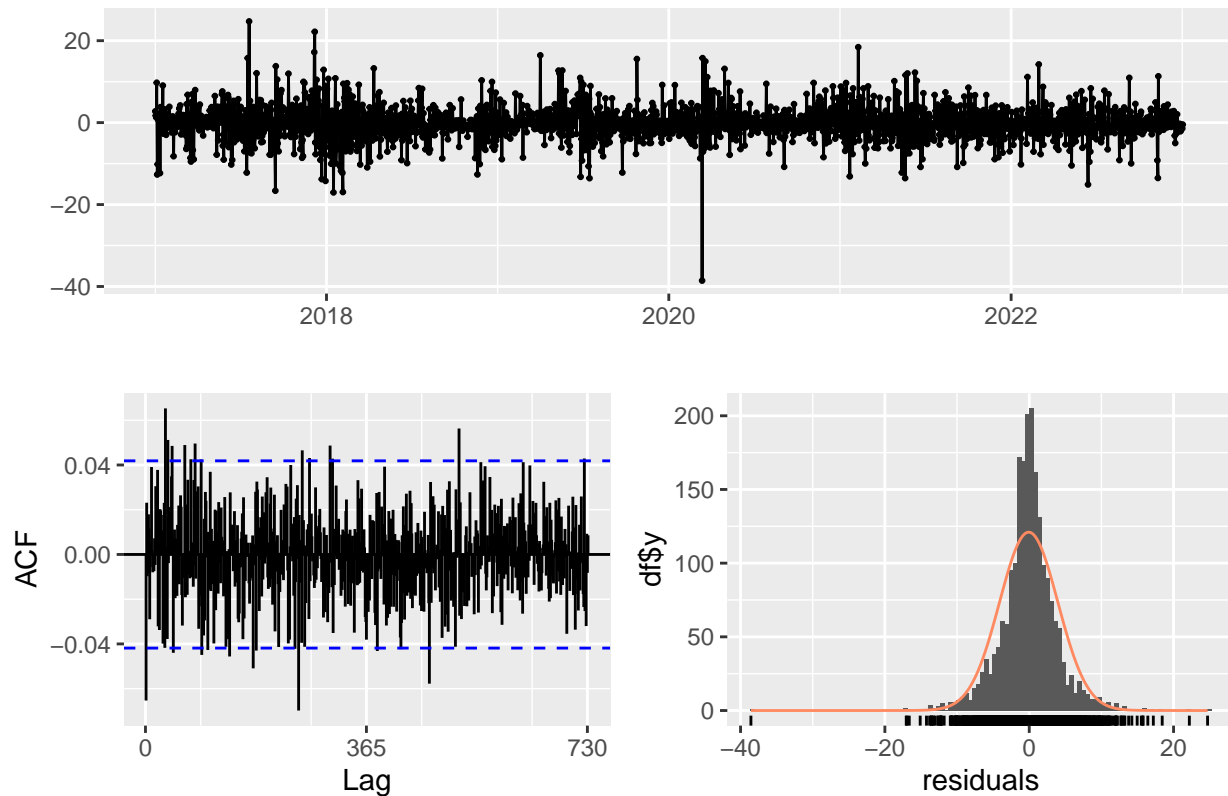


```
autoplot(ts_bitcoin_test) + autolayer(TBATS_forecast, series="TBATS", PI=FALSE)
```



```
TBATS_scores <- accuracy(TBATS_forecast$mean,msts_bitcoin_test)
#print(TBATS_scores)
checkresiduals(TBATS_fit)
```

## Residuals from TBATS




```
##
##  Ljung-Box test
##
## data:  Residuals from TBATS
## Q* = 449.5, df = 438, p-value = 0.3417
##
## Model df: 0.   Total lags used: 438
```

### 3.3 Arima with seasonality


```
#Arima with seasonality
arima_forecast <- forecast(auto.arima(msts_bitcoin,D=1),h=266)

plot(arima_forecast)
```

## Forecasts from ARIMA(2,0,0)(0,1,0)[365] with drift



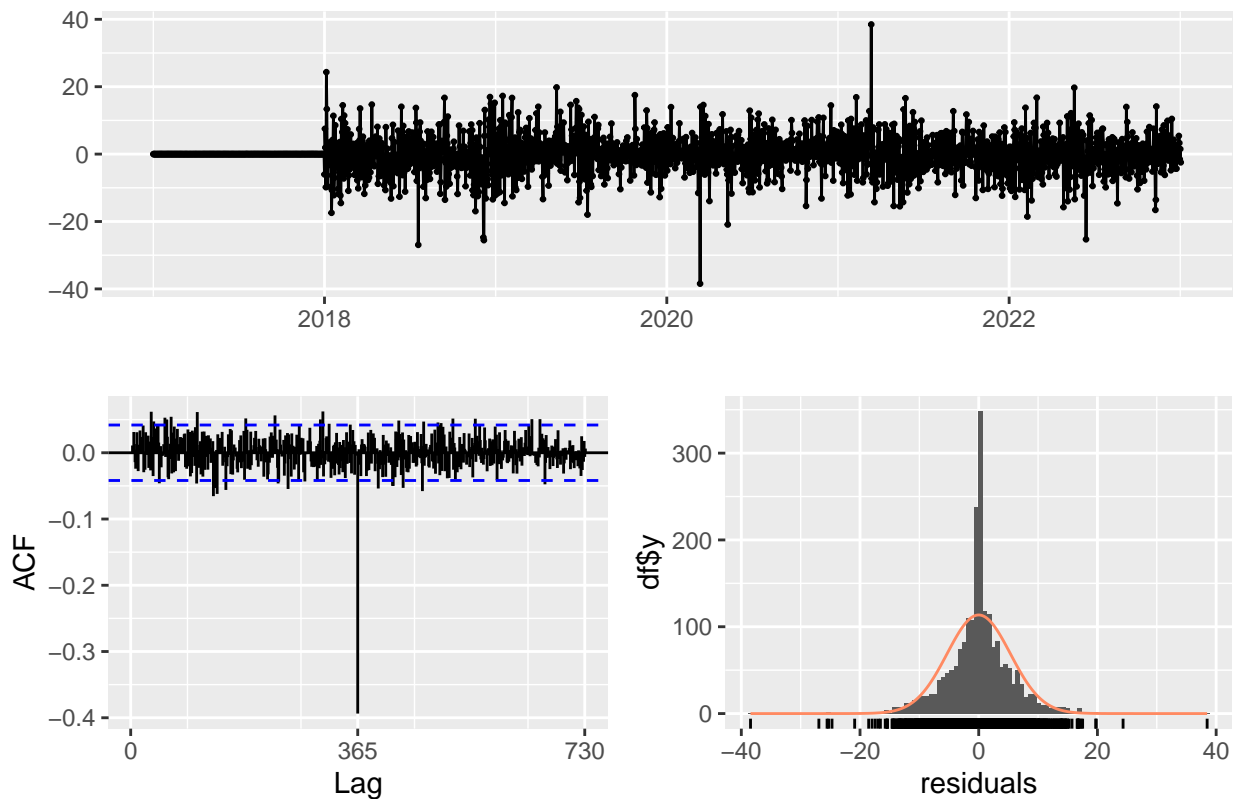
```
autoplot(msts_bitcoin_test) + autolayer(arima_forecast, series="ARIMA", PI=FALSE)
```



```
ARIMA_scores <- accuracy(arima_forecast$mean, msts_bitcoin_test)
#print(ARIMA_scores)
```

```
checkresiduals(arima_forecast)
```

Residuals from ARIMA(2,0,0)(0,1,0)[365] with drift




```
##  
## Ljung-Box test  
##  
## data: Residuals from ARIMA(2,0,0)(0,1,0)[365] with drift  
## Q* = 995.59, df = 436, p-value < 2.2e-16  
##  
## Model df: 2. Total lags used: 438
```


### 3.4 Neural network with fourier

```
#Neural network with fourier (3,12), p=1,P=1  
NN_fit1 <- nnetar(msts_bitcoin,p=1,P=1,xreg=fourier(msts_bitcoin, K=c(3,12)))  
NN_for1 <- forecast(NN_fit1,h=266, xreg=fourier(msts_bitcoin, K=c(3,12),h=266))  
  
plot(NN_for1)
```

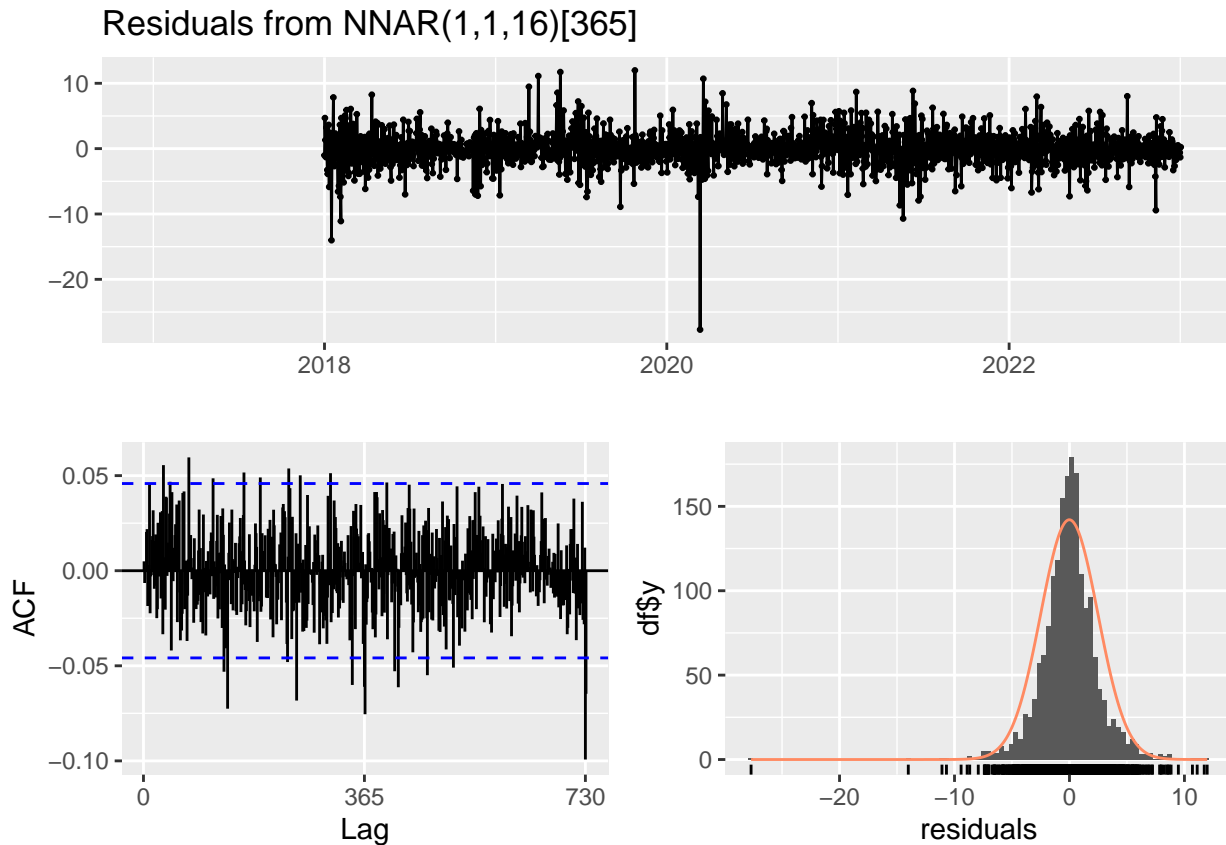
## Forecasts from NNAR(1,1,16)[365]



```
autoplot(msts_bitcoin_test) +  
  autolayer(NN_for1, series="Neural Network",PI=FALSE)
```



```
NN_scores1 <- accuracy(NN_for1$mean,msts_bitcoin_test)  
#print(NN_scores1)  
checkresiduals(NN_fit1)
```




```
##
##  Ljung-Box test
##
## data:  Residuals from NNAR(1,1,16)[365]
## Q* = 450.76, df = 438, p-value = 0.3266
##
## Model df: 0.   Total lags used: 438
```

#### 4. Models fit with external variables as regressor


```
# create the covid regressor
# Create a test series with 266 rows, each initialized to 0
msts_covid <- msts(train_bitcoin$covid,
                   seasonal.periods = c( 91.25, 365.25),
                   start = c(2017, 01, 01))
covid_regressors <- as.matrix(data.frame(fourier(msts_bitcoin, K=c(3,12), h=nrow(train_bitcoin)), "covid"))
future_covid_regressors <- as.matrix(data.frame(fourier(msts_bitcoin, K=c(3,12), h=nrow(test_bitcoin)), "covid"))

# NN+covid
NN_fit3 <- nnetar(msts_bitcoin, p=1, P=1, xreg=covid_regressors)
NN_for3 <- forecast(NN_fit3, h=266, xreg=future_covid_regressors)
plot(NN_for3)
```

## Forecasts from NNAR(1,1,17)[365]




```
autoplot(msts_bitcoin_test) +  
  autolayer(NN_for3, series="NN+covid",PI=FALSE)
```



```
NN_scores3 <- accuracy(NN_for3$mean,msts_bitcoin_test)  
#print(NN_scores3)  
checkresiduals(NN_fit3)
```


Residuals from NNAR(1,1,17)[365]




```
##
##  Ljung-Box test
##
## data:  Residuals from NNAR(1,1,17)[365]
## Q* = 438, df = 438, p-value = 0.491
##
## Model df: 0.   Total lags used: 438
## SS Exponential smoothing
SSES_fit1 <- es(msts_bitcoin,model="ZZZ",h=266,holdout=FALSE)
## Warning: Only additive models are allowed for your data. Changing the selection
## mechanism.
plot(SSES_fit1)
```


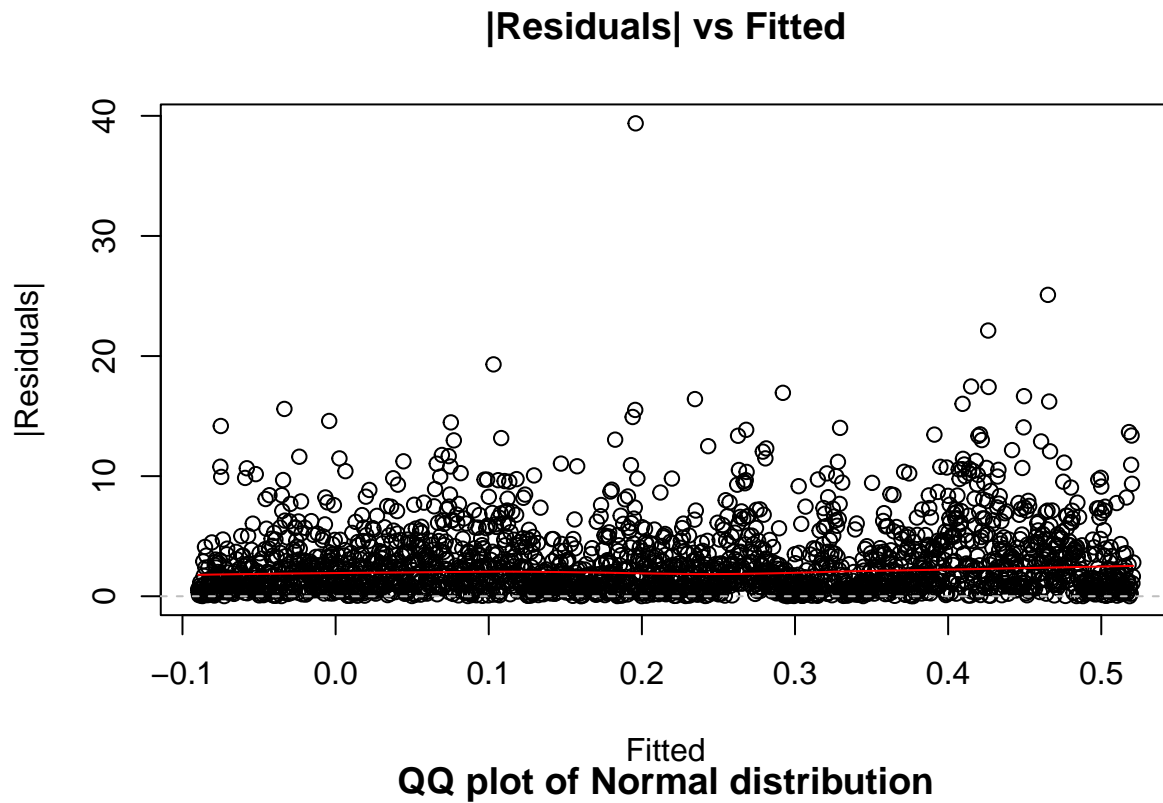


## Actuals vs Fitted



### Fitted Standardised Residuals vs Fitted






```
SSES_scores1 <- accuracy(SSES_fit1$forecast,msts_bitcoin_test)
#print(SSES_scores1)
```

```
#Arima+covid as regressor
```


```
ARIMA_fit1<-auto.arima(msts_bitcoin,seasonal= FALSE, lambda=1,xreg=covid_regressors)
```

```
ARIMA_for1<-forecast(ARIMA_fit1,xreg=future_covid_regressors,h=266)
plot(ARIMA_for1)
```

### Forecasts from Regression with ARIMA(5,1,0) errors

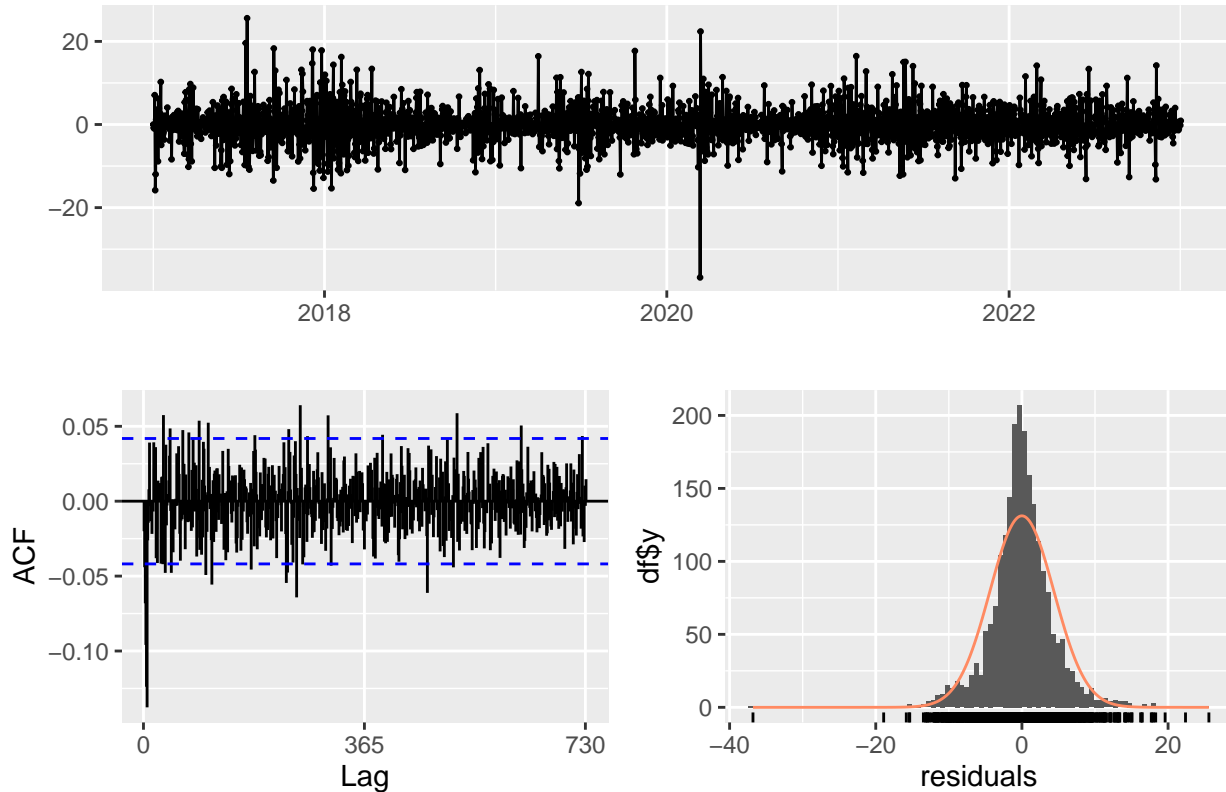


```
autoplot(msts_bitcoin_test) +
  autolayer(ARIMA_for1, series="ARIMA+Covid",PI=FALSE)
```



```
ARIMA_scores1 <- accuracy(ARIMA_for1$mean,msts_bitcoin_test)
#print(ARIMA_scores1)
checkresiduals(ARIMA_fit1)
```

### Residuals from Regression with ARIMA(5,1,0) errors



```
##
## Ljung-Box test
##
## data: Residuals from Regression with ARIMA(5,1,0) errors
## Q* = 597.87, df = 433, p-value = 2.343e-07
##
## Model df: 5. Total lags used: 438

#Crude oil regressors
oil_regressors<- as.matrix(data.frame(fourier(msts_bitcoin,K=c(3,12)),"oil"=msts_oil))
oil_fc<-forecast(msts_oil, h=266)
oil_regressors_fc<-as.matrix(data.frame(fourier(msts_bitcoin,K=c(3,12), h=266),"oil"= oil_fc$mean))

#Arima with regressor

arima_fit_r <- auto.arima(msts_bitcoin,seasonal= FALSE, lambda=0,xreg=oil_regressors)


## Warning in log(x): NaNs produced

## Warning in log(x): NaNs produced


arima_forecast_r <- forecast(arima_fit_r, xreg=oil_regressors_fc,h=266)

plot(arima_forecast_r)
```

## Forecasts from Regression with ARIMA(3,1,0) errors




```
autoplot(msts_bitcoin_test) + autolayer(arima_forecast_r, series="ARIMA+Oil", PI=FALSE)
```




```
autoplot(arima_forecast_r)
```

Forecasts from Regression with ARIMA(3,1,0) errors






```
autoplot(msts_bitcoin) +  
  autolayer(arima_forecast_r, series="ARIMA+Oil",PI=FALSE)
```



```
ARIMA_scores_oil <- accuracy(arima_forecast_r$mean,msts_bitcoin_test)
#print(ARIMA_scores_oil)

# Use checkresiduals to plot and assess residuals
checkresiduals(arima_forecast_r)
```


Residuals from Regression with ARIMA(3,1,0) errors




```
##
##  Ljung-Box test
##
## data:  Residuals from Regression with ARIMA(3,1,0) errors
## Q* = 1507.3, df = 435, p-value < 2.2e-16
##
## Model df: 3.   Total lags used: 438
NN_fit_o <- nnetar(msts_bitcoin,p=1,P=1,xreg=oil_regressors)
NN_fc_o <- forecast(NN_fit_o,h=266, xreg=oil_regressors_fc)

plot(NN_fc_o)
```

## Forecasts from NNAR(1,1,17)[365]




```
autoplot(msts_bitcoin_test) + autolayer(NN_fc_o, series="Neural Network+Oil", PI=FALSE)
```




```
autoplot(NN_fc_o)
```



Forecasts from NNAR(1,1,17)[365]



```
autoplot(msts_bitcoin) +  
  autolayer(NN_fc_o, series="Neural Network+Oil",PI=FALSE)
```




```

NN_scores_o <- accuracy(NN_fc_o$mean,msts_bitcoin_test)
#print(NN_scores_o)

# Use checkresiduals to plot and assess residuals
checkresiduals(NN_fc_o)

```

Residuals from NNAR(1,1,17)[365]



```

##
## Ljung-Box test
##
## data: Residuals from NNAR(1,1,17)[365]
## Q* = 490.47, df = 438, p-value = 0.04192
##
## Model df: 0. Total lags used: 438

```

```

# Combine two regressors

```

```

covid_oil_regressor <- as.matrix(data.frame(fourier(msts_bitcoin, K=c(3,12), h=nrow(train_bitcoin)), "c
covid_oil_regressor_fc <- as.matrix(data.frame(fourier(msts_bitcoin, K=c(3,12),h=nrow(test_bitcoin)), "

```

```

#Arima with two regressor

```

```

arima_fit_co <- auto.arima(msts_bitcoin,seasonal= FALSE, lambda=0,xreg=covid_oil_regressor)

```

```

## Warning in log(x): NaNs produced

```

```

## Warning in log(x): NaNs produced

```


```

arima_forecast_co <- forecast(arima_fit_co, xreg=covid_oil_regressor_fc,h=266)


```

```
plot(arima_forecast_co)
```

## Forecasts from Regression with ARIMA(0,0,0) errors




```
autoplot(msts_bitcoin_test) + autolayer(arima_forecast_co, series="ARIMA+Covid+Oil", PI=FALSE)
```




```
autoplot(arima_forecast_co)
```

Forecasts from Regression with ARIMA(0,0,0) errors

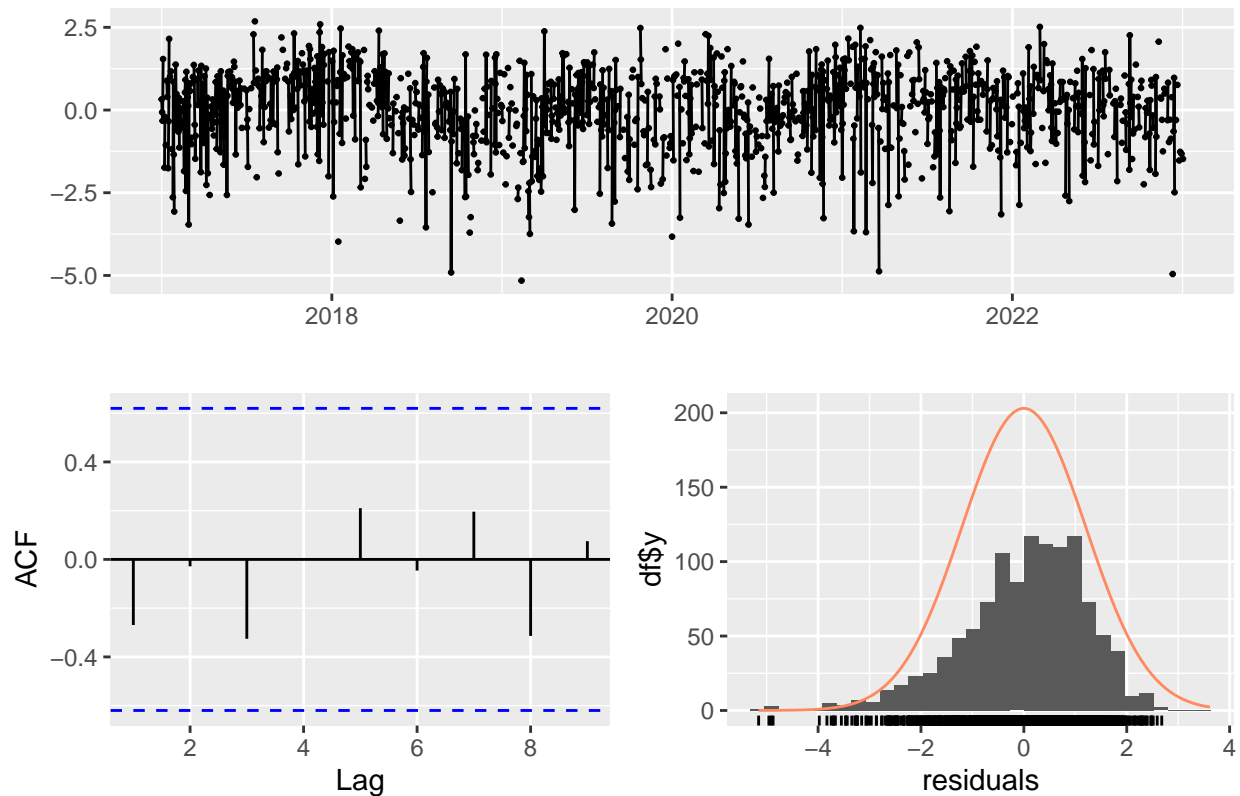


```
autoplot(msts_bitcoin) +  
  autolayer(arima_forecast_co, series="ARIMA+Covid+Oil",PI=FALSE)
```



```
ARIMA_scores_co <- accuracy(arima_forecast_co$mean,msts_bitcoin_test)
#print(ARIMA_scores_co)
# Use checkresiduals to plot and assess residuals
checkresiduals(arima_forecast_co)
```


Residuals from Regression with ARIMA(0,0,0) errors




```
##
##  Ljung-Box test
##
## data:  Residuals from Regression with ARIMA(0,0,0) errors
## Q* = 3678.2, df = 438, p-value < 2.2e-16
##
## Model df: 0.   Total lags used: 438
# NN+covid+oil
NN_fit_co <- nnetar(msts_bitcoin,p=1,P=1,xreg=covid_oil_regressor)
NN_for_co <- forecast(NN_fit_co,h=266, xreg=covid_oil_regressor_fc)

plot(NN_for_co)
```

## Forecasts from NNAR(1,1,18)[365]




```
autoplot(msts_bitcoin_test) + autolayer(NN_for_co, series="Neural Network+Covid+Oil", PI=FALSE)
```




```
autoplot(NN_for_co)
```

Forecasts from NNAR(1,1,18)[365]



```
autoplot(msts_bitcoin) +  
  autolayer(NN_for_co, series="NN+covid+oil",PI=FALSE)
```



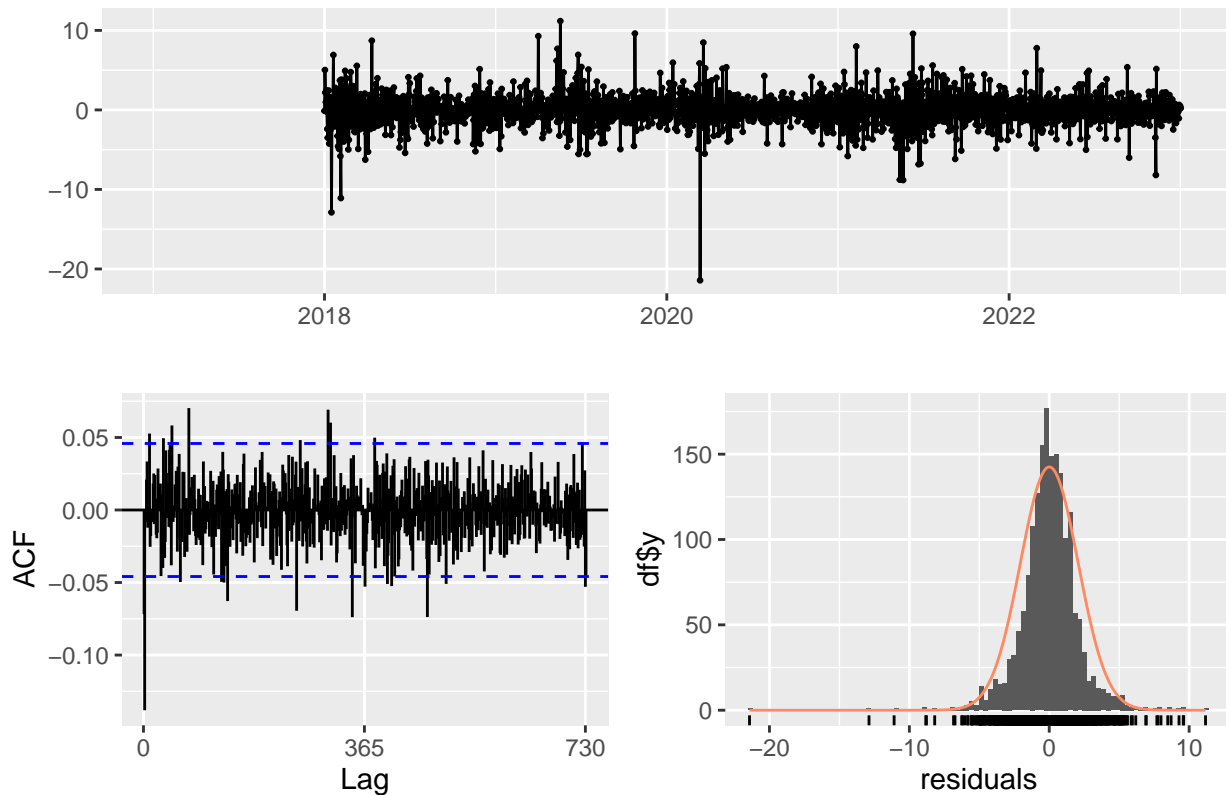


```

NN_scores_co <- accuracy(NN_for_co$mean,msts_bitcoin_test)
#print(NN_scores_co)
# Use checkresiduals to plot and assess residuals
checkresiduals(NN_for_co)

```

### Residuals from NNAR(1,1,18)[365]



```

##
##  Ljung-Box test
##
## data:  Residuals from NNAR(1,1,18)[365]
## Q* = 483.98, df = 438, p-value = 0.06367
##
## Model df: 0.   Total lags used: 438

```

## 5. Forecast

```

ts_oil_all <- final_data_filled %>% filter(
  DATE >= ymd("2017-01-01"))
combined_msts <- msts(
  ts_oil_all$DCOILBRENTU,
  seasonal.periods = c( 91.25,365.25),
  start=c(2017,01,01))
ts_bit_all <- bitcoin %>% filter(
  date >= ymd("2017-01-01"))
combined_msts_bit <- msts(
  ts_bit_all$Change,
  seasonal.periods = c( 91.25,365.25),
  start=c(2017,01,01))
oil_regressors_pd<- as.matrix(
  data.frame(fourier(combined_msts_bit,K=c(3,12)),
  "oil"=combined_msts))
oil_fc_pd<-forecast(combined_msts, h=30)
oil_regressors_fc_pd<-as.matrix(
  data.frame(fourier(combined_msts_bit,K=c(3,12), h=30),
  "oil"= oil_fc_pd$

```

```

arima_fit_r_pd <- auto.arima(combined_msts_bit,seasonal= FALSE, lambda=0,xreg=oil_regressors_pd)


## Warning in log(x): NaNs produced

## Warning in log(x): NaNs produced
arima_forecast_r_pd <- forecast(arima_fit_r_pd, xreg=oil_regressors_fc_pd,h=30)

plot(arima_forecast_r_pd)

```


## Forecasts from Regression with ARIMA(4,1,0) errors



```


autoplot(combined_msts_bit) + autolayer(arima_forecast_r_pd, series="ARIMA+Oil", PI=TRUE)

```




```
autoplot(arima_forecast_r_pd)
```

Forecasts from Regression with ARIMA(4,1,0) errors




```
autoplot(combined_msts_bit) +  
  autolayer(arima_forecast_r_pd, series="ARIMA+Oil",PI=FALSE)
```



```
# Extract the mean forecast values  
mean_forecast <- arima_forecast_r_pd$mean  
  
# Length of mean_forecast  
n <- length(mean_forecast)  
  
# Extract the dates for the x-axis  
dates <- seq(from = as.Date("2024-04-08"), by = "day", length.out = n)  
  
# Plot the mean forecast with specified x-axis range  
plot(dates, mean_forecast, type = "l", xlab = "Date", ylab = "Forecasted Values",  
     main = "ARIMA Forecast", xlim = c(as.Date("2024-04-08"), as.Date("2024-05-08")))
```

## ARIMA Forecast



## 6.

Table

```
# Load required packages
```

```
library(knitr)
```

```
library(kableExtra)
```

```
##
```

```
## Attaching package: 'kableExtra'
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
## group_rows
```

```
# Combine scores
```

```
combined_scores <- rbind(
```

```
  ETS_scores,
```

```
  TBATS_scores,
```

```
  ARIMA_scores,
```

```
  ARIMA_scores1,
```

```
  NN_scores1,
```

```
  NN_scores3,
```

```
  ARIMA_scores_oil,
```

```
  ARIMA_scores_co,
```

```
  NN_scores_o,
```

```
  NN_scores_co
```

```
)
```

```
# Define row names
```

```
rownames(combined_scores) <- c(
```

```
  "ETS", "TBATS", "ARIMA", "ARIMA with Covid",
```

```
  "Neural Network with Fourier", "Neural Network with Covid",
```

```
  "ARIMA with Oil", "ARIMA with Covid and Oil",
```

```

    "Neural Network with Oil", "Neural Network with Covid and Oil"
)

knitr::kable(
  combined_scores,
  format = "latex",
  caption = "Combined Scores"
) %>%
  kable_styling(
    latex_options = c("hold_position", "scale_down"),
    full_width = FALSE
  )

```

Table 1: Combined Scores

	ME	RMSE	MAE	MPE	MAPE	ACF1	Theil's U
ETS	0.6482149	3.360422	2.599518	86.34347	498.0130	-0.1136271	1.283959
TBATS	0.5822264	2.701150	1.852876	99.08521	205.2798	0.0192255	1.078225
ARIMA	0.8172699	4.580938	3.314904	414.25838	1024.7176	0.0374629	1.634431
ARIMA with Covid	1.5937054	3.013271	2.235520	61.13900	464.8829	-0.0445094	1.155884
Neural Network with Fourier	0.3654727	3.079558	2.242687	185.61673	414.9941	0.0152402	1.037852
Neural Network with Covid	0.4074562	2.823624	2.128092	114.48023	445.0392	-0.0417718	1.035712
ARIMA with Oil	-0.1168176	2.511553	1.699530	95.62725	177.5840	-0.0847061	1.115706
ARIMA with Covid and Oil	-1.2261912	2.803823	2.148202	100.87372	472.3427	-0.0731229	1.658998
Neural Network with Oil	0.3001861	2.964914	2.200633	205.13294	403.2813	0.0648576	1.538607
Neural Network with Covid and Oil	0.3705762	2.995247	2.197752	246.33732	412.4288	0.0413739	1.300681